

1. Concurrency Models (Pros & Cons)

1. Threads:

- **Pros:** Great for CPU-heavy tasks and can run tasks in parallel on multiple cores.
- **Cons:** Risk of issues like race conditions, higher complexity, and context-switching overhead.

2. Event-Driven:

- **Pros:** Efficient for I/O-bound tasks, lightweight, often used in web servers for low latency.
- **Cons:** Harder to manage as it requires handling states, not ideal for CPU-bound work.

3. Coroutines:

- **Pros:** Efficient for single-threaded, I/O-heavy applications with lower memory overhead.
- **Cons:** Limited to single-threaded execution, so it's not suitable for CPU-bound tasks.

2. Concurrency vs Parallelism

- **Concurrency:** Manages multiple tasks by switching between them, not necessarily at the same time. Think of it as task coordination to maximize resource use.
- **Parallelism:** Executes tasks simultaneously, often on multiple processors. It's about speed-up, dividing tasks to run at the same time for CPU-bound work.

In Short: Concurrency is task management; parallelism is task execution.

3. Blocking vs Non-Blocking Concurrency Algorithms

1. Blocking:

- **Usage:** A task waits until a resource is available, which is simpler but can cause delays.
- **Example:** Traditional file I/O where tasks wait until resources are released.
- **Drawback:** Reduces performance, prone to deadlocks.

2. Non-Blocking:

- **Usage:** Tasks don't wait; they try again later or use a different approach if resources aren't available.
- **Example:** Async I/O or lock-free algorithms that retry when resources are unavailable.
- **Drawback:** More complex to implement but ideal for high-performance, scalable systems.